

Mercedes-Benz Greener Manufacturing

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
[2]: train = pd.read_csv('train.csv', index_col = 'ID')
test = pd.read_csv('test.csv', index_col = 'ID')
```

```
[3]: train.head()
```

```
[3]:      y  X0 X1  X2 X3 X4 X5 X6 X8  X10  ...  X375  X376  X377  X378  X379  \
ID
0  130.81  k  v  at  a  d  u  j  o    0  ...    0    0    1    0    0
6   88.53  k  t  av  e  d  y  l  o    0  ...    1    0    0    0    0
7   76.26  az  w   n  c  d  x  j  x    0  ...    0    0    0    0    0
9   80.62  az  t   n  f  d  x  l  e    0  ...    0    0    0    0    0
13  78.02  az  v   n  f  d  h  d  n    0  ...    0    0    0    0    0
```

```
      X380  X382  X383  X384  X385
ID
0         0     0     0     0     0
6         0     0     0     0     0
7         0     1     0     0     0
9         0     0     0     0     0
13        0     0     0     0     0
```

[5 rows x 377 columns]

Replacing strings with numbers in train and test dataframes. Note that a combined list of all unique strings is prepared for each feature (containing string) for both train and test data before replacing it with numbers. This is done to ensure that each strings gets mapped to same number for both train and test data.

```
[4]: for col in train.columns:
      if(train[col].dtype != np.float64 and train[col].dtype != np.int64):

          # making a list of unique strings in train and test feature
```

```

unique_train = train[col].unique().tolist()
unique_test = test[col].unique().tolist()

# making a combined list
for member in unique_test:
    if member not in unique_train:
        unique_train.append(member)

# mapping with numbers
map_dict = dict(zip(unique_train, range(len(unique_train))))
train[col] = train[col].replace(to_replace = map_dict)
test[col] = test[col].replace(to_replace = map_dict)

```

```
[5]: train.head()
```

```

[5]:      y  X0  X1  X2  X3  X4  X5  X6  X8  X10  ...  X375  X376  X377  X378  \
ID
0  130.81  0  0  0  0  0  0  0  0  0  ...    0    0    1    0
6   88.53  0  1  1  1  0  1  1  0  0  ...    1    0    0    0
7   76.26  1  2  2  2  0  2  0  1  0  ...    0    0    0    0
9   80.62  1  1  2  3  0  2  1  2  0  ...    0    0    0    0
13  78.02  1  0  2  3  0  3  2  3  0  ...    0    0    0    0

```

```

      X379  X380  X382  X383  X384  X385
ID
0         0     0     0     0     0     0
6         0     0     0     0     0     0
7         0     0     1     0     0     0
9         0     0     0     0     0     0
13        0     0     0     0     0     0

```

[5 rows x 377 columns]

Checking if train or test data has any NaN value. Also, getting summary of data

```

[6]: print(train.isnull().values.any())
      print(test.isnull().values.any())
      train.describe()

```

False

False

```

[6]:      count      y      X0      X1      X2      X3  \
count  4209.000000  4209.000000  4209.000000  4209.000000  4209.000000
mean    100.669318    12.110715    6.467569    7.851509    2.415301
std     12.679381     8.315637    4.789927    5.644031    1.361654
min      72.110000     0.000000    0.000000    0.000000    0.000000

```

25%	90.820000	6.000000	3.000000	4.000000	2.000000
50%	99.150000	11.000000	6.000000	7.000000	2.000000
75%	109.010000	15.000000	7.000000	10.000000	3.000000
max	265.320000	46.000000	26.000000	43.000000	6.000000

	X4	X5	X6	X8	X10	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.002138	16.839629	3.031124	11.457591	0.013305
std	0.073900	6.357474	2.554581	7.040194	0.114590
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	11.000000	1.000000	5.000000	0.000000
50%	0.000000	17.000000	2.000000	12.000000	0.000000
75%	0.000000	22.000000	6.000000	17.000000	0.000000
max	3.000000	28.000000	11.000000	24.000000	1.000000

	X375	X376	X377	X378	X379	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.318841	0.057258	0.314802	0.020670	0.009503
std	0.466082	0.232363	0.464492	0.142294	0.097033
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	1.000000	0.000000	1.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

	X380	X382	X383	X384	X385	...	\
count	4209.000000	4209.000000	4209.000000	4209.000000	4209.000000
mean	0.008078	0.007603	0.001663	0.000475	0.001426
std	0.089524	0.086872	0.040752	0.021796	0.037734
min	0.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	0.000000	0.000000	0.000000	0.000000
50%	0.000000	0.000000	0.000000	0.000000	0.000000
75%	0.000000	0.000000	0.000000	0.000000	0.000000
max	1.000000	1.000000	1.000000	1.000000	1.000000

[8 rows x 377 columns]

Splitting features and labels. Also, performing min_max scaling on features

```
[7]: X_train = train.iloc[:,1:]
      y_train = train.iloc[:,0]

      from sklearn.preprocessing import MinMaxScaler
      scaling = MinMaxScaler().fit(X_train)
      X_train_scaled = scaling.transform(X_train)
      test_scaled = scaling.transform(test)
```

Regression with linear model. Cross-validation score shows that the linear model

performs very poorly.

```
[8]: from sklearn.linear_model import LinearRegression
      from sklearn.model_selection import cross_val_score

      reg = LinearRegression()
      print(cross_val_score(reg, X_train_scaled, y_train, cv=10))
```

```
[-1.33042147e+23 -1.91338315e+22 -4.94938240e+22 -1.33714229e+24
 -9.64726715e+22 -2.99456830e+21 -4.23383313e+23 -1.81628722e+23
 -6.40148587e+22 -7.38634788e+23]
```

Regression with Lasso model (L1 regularization). As the number of features are very large, Lasso regularization would assign lesser weights to non-important features and in-turn reduce their contribution in the final regression model.

```
[9]: from sklearn.linear_model import Lasso
      from sklearn.model_selection import GridSearchCV

      grid_values = {'alpha': [0.0235, 0.024, 0.0245]}
      grid_lasso = GridSearchCV(Lasso(), param_grid = grid_values, cv=10, scoring =_
        ↪ 'r2')
      grid_lasso.fit(X_train_scaled, y_train)
      predict_lasso = grid_lasso.predict(test_scaled)

      print('Mean score matrix: ', grid_lasso.cv_results_['mean_test_score'])
      print('Grid best parameter (max. accuracy): ', grid_lasso.best_params_)
      print('Grid best score (accuracy): ', grid_lasso.best_score_)
```

```
Mean score matrix: [0.56300014 0.56300502 0.56299962]
Grid best parameter (max. accuracy): {'alpha': 0.024}
Grid best score (accuracy): 0.5630050235597193
```

Let's also try Ridge regression (L2 regularization)

```
[15]: from sklearn.linear_model import Ridge
      from sklearn.model_selection import GridSearchCV

      grid_values = {'alpha': [40, 40.5, 41]}
      grid_ridge = GridSearchCV(Ridge(), param_grid = grid_values, cv=10, scoring =_
        ↪ 'r2')
      grid_ridge.fit(X_train_scaled, y_train)
      predict_ridge = grid_ridge.predict(test_scaled)

      print('Mean score matrix: ', grid_ridge.cv_results_['mean_test_score'])
      print('Grid best parameter (max. accuracy): ', grid_ridge.best_params_)
      print('Grid best score (accuracy): ', grid_ridge.best_score_)
```

```
Mean score matrix: [0.55378705 0.55378761 0.55378743]
```

Grid best parameter (max. accuracy): {'alpha': 40.5}
Grid best score (accuracy): 0.5537876100313096

Regression with Xgboost. It shows the best R2 score. We will use this model to do final predictions.

```
[16]: pip install xgboost
```

```
Requirement already satisfied: xgboost in c:\users\91805\anaconda1\lib\site-  
packages (1.7.5)  
Requirement already satisfied: scipy in c:\users\91805\anaconda1\lib\site-  
packages (from xgboost) (1.9.1)  
Requirement already satisfied: numpy in c:\users\91805\anaconda1\lib\site-  
packages (from xgboost) (1.21.5)  
Note: you may need to restart the kernel to use updated packages.
```

```
[17]: import xgboost as xgb
```

```
[18]: conda install -c conda-forge xgboost
```

```
Collecting package metadata (current_repodata.json): ...working... done  
Solving environment: ...working... done
```

```
# All requested packages already installed.
```

Note: you may need to restart the kernel to use updated packages.

```
[ ]: grid_values = {'n_estimators': [100,105,106], 'learning_rate': [0.13,0.135,0.  
    ↪14], 'max_depth': [1,2,3]}  
grid_xgb = GridSearchCV(xgb.XGBRegressor(), param_grid = grid_values, cv=10,   
    ↪scoring = 'r2')  
grid_xgb.fit(X_train_scaled, y_train)  
predict_xgb = grid_xgb.predict(test_scaled)  
  
print('Mean score matrix: ', grid_xgb.cv_results_['mean_test_score'])  
print('Grid best parameter (max. accuracy): ', grid_xgb.best_params_)  
print('Grid best score (accuracy): ', grid_xgb.best_score_)
```

```
Mean score matrix: [ 0.55597401 0.55620147 0.55675942 0.58165304 0.5818088 0.58176478  
0.57938243 0.57941404 0.57938819 0.55721329 0.55745941 0.55774666 0.58204994 0.58214103  
0.58202328 0.57727201 0.57720783 0.57713946 0.55807455 0.55816391 0.55836724 0.58153668  
0.58149164 0.58166668 0.57805699 0.57808365 0.57817218]
```

```
Grid best parameter (max. accuracy): {'learning_rate': 0.135, 'max_depth': 2, 'n_estimators':  
75}
```

```
Grid best score (accuracy): 0.582141029072
```

```
[ ]: final_predictions = pd.DataFrame()
      final_predictions['id'] = test.index
      final_predictions['y'] = pd.Series(predict_xgb)
      final_predictions.to_csv('predictions.csv', index=False)
```